

In [83]: `#pip install nbconvert`

In [2]: `# Import Libraries`

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

In [3]: `pd.set_option('display.max_rows',None)`  
`pd.set_option('display.max_columns',None)`

## Reading Mercedes Benz Data Set

In [4]: `test = pd.read_csv('E:/downloads/MACHINE LEARNING PROJECTS/Mercedes-Benz Greener`  
`train =pd.read_csv('E:/downloads/MACHINE LEARNING PROJECTS/Mercedes-Benz Greener`

In [5]: `test.head()`

Out[5]:

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	X11	X12	X13	X14	X15	X16	X17	X18	X19	X2
0	1	az	v	n	f	d	t	a	w	0	0	0	0	0	0	0	0	0	0	0
1	2	t	b	ai	a	d	b	g	y	0	0	0	0	0	0	0	0	0	0	1
2	3	az	v	as	f	d	a	j	j	0	0	0	0	1	0	0	0	0	0	0
3	4	az	l	n	f	d	z	l	n	0	0	0	0	0	0	0	0	0	0	0
4	5	w	s	as	c	d	y	i	m	0	0	0	0	1	0	0	0	0	0	0

In [6]: `train.head()`

Out[6]:

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	X10	X11	X12	X13	X14	X15	X16	X17	X18
0	0	130.81	k	v	at	a	d	u	j	o	0	0	0	1	0	0	0	0	1
1	6	88.53	k	t	av	e	d	y	l	o	0	0	0	0	0	0	0	0	1
2	7	76.26	az	w	n	c	d	x	j	x	0	0	0	0	0	0	0	1	0
3	9	80.62	az	t	n	f	d	x	l	e	0	0	0	0	0	0	0	0	0
4	13	78.02	az	v	n	f	d	h	d	n	0	0	0	0	0	0	0	0	0

In [7]: `train.shape`

Out[7]: (4209, 378)

In [8]: `train.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Columns: 378 entries, ID to X385
dtypes: float64(1), int64(369), object(8)
memory usage: 12.1+ MB
```

In [9]: `test.shape`

Out[9]: (4209, 377)

In [10]: `test.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Columns: 377 entries, ID to X385
dtypes: int64(369), object(8)
memory usage: 12.1+ MB
```

In [11]: `train.describe()`

Out[11]:

	ID	y	X10	X11	X12	X13	X14	
<b>count</b>	4209.000000	4209.000000	4209.000000	4209.0	4209.000000	4209.000000	4209.000000	4209.000000
<b>mean</b>	4205.960798	100.669318	0.013305	0.0	0.075077	0.057971	0.428130	0.428130
<b>std</b>	2437.608688	12.679381	0.114590	0.0	0.263547	0.233716	0.494867	0.494867
<b>min</b>	0.000000	72.110000	0.000000	0.0	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	2095.000000	90.820000	0.000000	0.0	0.000000	0.000000	0.000000	0.000000
<b>50%</b>	4220.000000	99.150000	0.000000	0.0	0.000000	0.000000	0.000000	0.000000
<b>75%</b>	6314.000000	109.010000	0.000000	0.0	0.000000	0.000000	1.000000	1.000000
<b>max</b>	8417.000000	265.320000	1.000000	0.0	1.000000	1.000000	1.000000	1.000000

In [12]: `train.nunique()`

Out[12]:

ID	4209
y	2545
X0	47
X1	27
X2	44
X3	7
X4	4
X5	29
X6	12
X8	25
X10	2
X11	1
X12	2
X13	2
X14	2
X15	2
X16	2
X17	2
X18	2
X19	2

In [13]: `train['y'].nunique()`

Out[13]: 2545

In [14]: `print("y is the difference in train test columns, so it is a Target column")`  
`print("Since Target Variable has 2500 unique values, it is a Regression Problem.")`

y is the difference in train test columns, so it is a Target column  
 Since Target Variable has 2500 unique values, it is a Regression Problem.

In [15]: `#check the unique value , for that if unqiues value is 1 than it has no variance`  
`len(train['X0'].unique())`

Out[15]: 47

In [16]: `#find all the columns of our data`  
`col = train.columns`  
`col`

Out[16]: Index(['ID', 'y', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8',  
 ...,  
 'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X384',  
 'X385'],  
 dtype='object', length=378)

## Checking columns with zero variance

```
In [17]: #To check uniques value in all 384 columns seprately is tuff so we do it by a func
for each in col:
    if len(train[each].unique()) == 1:
        print('column name', ' ', each)
```

```
column name    X11
column name    X93
column name    X107
column name    X233
column name    X235
column name    X268
column name    X289
column name    X290
column name    X293
column name    X297
column name    X330
column name    X347
```

```
In [18]: #from above we have found that columns X11,X93,X107,X233,X235,X268,X289,X290,X293
#than from probelms objective -If for any column(s), the variance is equal to zero
#we remove
train.drop(['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X297', 'X330', 'X347'])
```

```
In [19]: train.shape
```

```
Out[19]: (4209, 366)
```

## Checking Null Value in Test & Train Data

```
In [20]: train.isnull().sum()
```

```
Out[20]: ID      0
y      0
X0      0
X1      0
X2      0
X3      0
X4      0
X5      0
X6      0
X8      0
X10     0
X12     0
X13     0
X14     0
X15     0
X16     0
X17     0
X18     0
X19     0
X20     0
```

```
In [21]: train.isnull().sum().sum()
```

```
Out[21]: 0
```

```
In [22]: train.isnull().sum().any()
```

```
Out[22]: False
```

```
In [23]: test.isnull().sum()
```

```
Out[23]: ID      0
          X0      0
          X1      0
          X2      0
          X3      0
          X4      0
          X5      0
          X6      0
          X8      0
          X10     0
          X11     0
          X12     0
          X13     0
          X14     0
          X15     0
          X16     0
          X17     0
          X18     0
          X19     0
          X20     0
```

```
In [24]: test.isnull().sum().sum()
```

```
Out[24]: 0
```

```
In [25]: test.isnull().sum().any()
```

```
Out[25]: False
```

## Checking columns data types

```
In [26]: train.dtypes.value_counts()
```

```
Out[26]: int64      357
          object      8
          float64     1
          dtype: int64
```

```
In [27]: #to find dtype of all columns we do like this
for each in train.columns:
    print(each, "---", np.dtype(train[each]))
```

```
ID --- int64
y --- float64
X0 --- object
X1 --- object
X2 --- object
X3 --- object
X4 --- object
X5 --- object
X6 --- object
X8 --- object
X10 --- int64
X12 --- int64
X13 --- int64
X14 --- int64
X15 --- int64
X16 --- int64
X17 --- int64
X18 --- int64
X19 --- int64
X20 --- int64
```

```
In [28]: category_cols = [c for c in train if train[c].dtype == np.dtype('object')]
category_cols
```

```
Out[28]: ['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8']
```

## Apply Label Encoder

```
In [29]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
```

```
In [30]: train[category_cols] = train[category_cols].apply(le.fit_transform)
```

```
In [31]: # Features
train_feature = train.drop(['ID', 'y'], axis = 1)
# Target
train_target = train['y']
```

In [32]: `train_feature.head()`

Out[32]:

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X12	X13	X14	X15	X16	X17	X18	X19	X20	X21	
0	32	23	17	0	3	24	9	14	0	0	1	0	0	0	0	1	0	0	0	1
1	32	21	19	4	3	28	11	14	0	0	0	0	0	0	0	1	0	0	0	0
2	20	24	34	2	3	27	9	23	0	0	0	0	0	0	1	0	0	0	0	0
3	20	21	34	5	3	27	11	4	0	0	0	0	0	0	0	0	0	0	0	0
4	20	23	34	5	3	12	3	13	0	0	0	0	0	0	0	0	0	0	0	0

In [33]: `train_target.head()`

Out[33]:

```
0    130.81
1     88.53
2     76.26
3     80.62
4     78.02
Name: y, dtype: float64
```

In [34]: `train_feature.shape`

Out[34]: (4209, 364)

In [35]: `train_target.shape`

Out[35]: (4209,)

In [36]: `train_feature.describe()`

Out[36]:

	X0	X1	X2	X3	X4	X5	X6
<b>count</b>	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000
<b>mean</b>	29.760751	11.113566	17.306486	2.919696	2.997862	13.340223	6.807318
<b>std</b>	13.738338	8.531001	10.899914	1.739912	0.073900	8.250832	2.916973
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	19.000000	3.000000	8.000000	2.000000	3.000000	5.000000	6.000000
<b>50%</b>	35.000000	13.000000	16.000000	2.000000	3.000000	15.000000	7.000000
<b>75%</b>	43.000000	20.000000	25.000000	5.000000	3.000000	21.000000	9.000000
<b>max</b>	46.000000	26.000000	43.000000	6.000000	3.000000	28.000000	11.000000

## Perfoming Dimensionality Reduction (PCA)

```
In [37]: print(train_feature.shape)
print(train_target.shape)
```

```
(4209, 364)
(4209,)
```

```
In [38]: from sklearn.decomposition import PCA
pca = PCA(n_components = 0.95)
```

```
In [39]: pca.fit(train_feature, train_target)
```

```
Out[39]: PCA(n_components=0.95)
```

```
In [40]: train_feature_pca = pca.fit_transform(train_feature)
```

```
In [41]: train_feature_pca.shape
```

```
Out[41]: (4209, 6)
```

## Splitting Data

```
In [42]: from sklearn.model_selection import train_test_split
```

```
In [43]: train_x, test_x, train_y, test_y = train_test_split(train_feature_pca, train_target,
```

```
In [44]: train_x.shape
```

```
Out[44]: (3367, 6)
```

```
In [45]: train_y.shape
```

```
Out[45]: (3367,)
```

```
In [46]: test_x.shape
```

```
Out[46]: (842, 6)
```

```
In [47]: test_y.shape
```

```
Out[47]: (842,)
```

## XGboost

```
In [48]: import xgboost as xgb
```

```
In [49]: xgb_reg = xgb.XGBRegressor(objective= 'reg:linear', colsample_bytree = 0.3, learni
```



```
In [50]: model = xgb_reg.fit(train_x,train_y)
```

```
[17:10:26] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.6.0/src/objective/regression_obj.cu:203: reg:linear is now deprecated in favor of reg:squarederror.
```

```
In [51]: y_pred = model.predict(test_x)
y_pre = model.predict(train_x)
```

```
In [52]: y_pred.shape
```

```
Out[52]: (842,)
```

```
In [53]: y_pre.shape
```

```
Out[53]: (3367,)
```

```
In [54]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [55]: print('Train_Score=',model.score(train_x, train_y))
print('Test_Score=',model.score(test_x, test_y))
```

```
Train_Score= 0.5701760216727423
Test_Score= 0.25325418172617364
```

```
In [56]: print(r2_score(test_y, y_pred))
print(r2_score(train_y, y_pre))
```

```
0.25325418172617364
0.5701760216727423
```

```
In [57]: print('MAE=', mean_absolute_error(test_y, y_pred))
print('MSE=',mean_squared_error(test_y, y_pred))
print('RSME=',np.sqrt(mean_squared_error(test_y, y_pred)))
```

```
MAE= 8.044838375354322
MSE= 121.97190419641204
RSME= 11.044089106685623
```

## Prediction on test data set



In [66]: `test_feature.head()`

Out[66]:

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X12	X13	X14	X15	X16	X17	X18	X19	X20	X21
0	az	v	n	f	d	t	a	w	0	0	0	0	0	0	0	0	0	0	0
1	t	b	ai	a	d	b	g	y	0	0	0	0	0	0	0	0	1	0	0
2	az	v	as	f	d	a	j	j	0	0	0	1	0	0	0	0	0	0	0
3	az	l	n	f	d	z	l	n	0	0	0	0	0	0	0	0	0	0	0
4	w	s	as	c	d	y	i	m	0	0	0	1	0	0	0	0	0	0	0

In [67]: `test_feature.describe(include='object')`

Out[67]:

	X0	X1	X2	X3	X4	X5	X6	X8
count	4209	4209	4209	4209	4209	4209	4209	4209
unique	49	27	45	7	4	32	12	25
top	ak	aa	as	c	d	v	g	e
freq	432	826	1658	1900	4203	246	1073	274

In [68]: `category_cols1 = [c for c in test_feature if test_feature[c].dtype == np.dtype('object')]
category_cols1`

Out[68]: `['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8']`

In [69]: `from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()`

In [70]: `test_feature[category_cols1] = test_feature[category_cols1].apply(le.fit_transform)`

In [71]: `test_feature.head()`

Out[71]:

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X12	X13	X14	X15	X16	X17	X18	X19	X20	X21
0	21	23	34	5	3	26	0	22	0	0	0	0	0	0	0	0	0	0	0
1	42	3	8	0	3	9	6	24	0	0	0	0	0	0	0	0	1	0	0
2	21	23	17	5	3	0	9	9	0	0	0	1	0	0	0	0	0	0	0
3	21	13	34	5	3	31	11	13	0	0	0	0	0	0	0	0	0	0	0
4	45	20	17	2	3	30	8	12	0	0	0	1	0	0	0	0	0	0	0

```
In [72]: test_feature.dtypes.value_counts()
```

```
Out[72]: int64    356  
         int32      8  
         dtype: int64
```

```
In [73]: pca.fit(test_feature)
```

```
Out[73]: PCA(n_components=0.95)
```

```
In [74]: test_feature_trans = pca.fit_transform(test_feature)
```

```
In [75]: test_feature_trans.shape
```

```
Out[75]: (4209, 6)
```

```
In [76]: test_feature_trans
```

```
Out[76]: array([[ 14.58336183,  14.16672593,  13.53857566,   2.40835691,  
                  11.31942221,   6.94220721],  
                [-15.25161267, -7.73675643, -7.45495068, -2.66203503,  
                  11.59379316,   1.15940345],  
                [ 11.8564649 , -1.68017324, -9.9896148 ,  14.91886587,  
                 -1.08886021, -2.69130553],  
                ...,  
                [-13.44644008,   3.2885825 , -6.85236431,  18.91025575,  
                  11.32365564,   3.22410016],  
                [ 24.92612317, -4.89888683, -10.16941028,  11.44337736,  
                   5.90178724,   4.55323232],  
                [-15.38430989, -7.73425491, -15.4930104 , -0.5595126 ,  
                   4.7793639 ,   1.0829113 ]])
```

```
In [77]: test_pred = model.predict(test_feature_trans)
```

```
In [78]: test_pred.shape
```

```
Out[78]: (4209,)
```

```
In [79]: test_pred
```

```
Out[79]: array([ 98.24064 ,  99.245674,  99.70339 , ...,  96.97598 , 111.26956 ,  
                 102.940125], dtype=float32)
```

```
In [80]: x = pd.DataFrame(test_pred)
```

In [81]: x

Out[81]:

	0
0	98.240639
1	99.245674
2	99.703392
3	103.333298
4	106.152199
5	93.206100
6	91.223396
7	98.224739
8	114.123512
9	104.349297
10	114.255074

**Thank you**